

p1

Sunday, February 28, 2021 4:54 PM

Architecture:

1 Hidden layer (256 units), Sigmoid Activation
 Output Layer (10 neurons): Softmax Cross Entropy Loss

hiddens = [256]
 lr = 0.05
 num_epochs = 15
 batch_size = 8

Test Error: 0.0794**Reflection:**

I noticed that using a smaller batch size like 8 may be slower to train since it needs to perform more weight updates per epoch, and the updates are also noisier (hence the 'zig zag' shapes in the graph). However, the validation loss decreases faster and converges towards the minimum faster.

- I decided to use 1 Hidden layer with sigmoid activation, and softmax cross entropy loss to classify the input into K=10 output classes. I experimented with 2 Hidden layers with sigmoid activation, but found that it does not achieve a significantly lower test error, and takes more time to train.

When observing the graph, it seems that the graph overfits when the training error continues to decrease, but the validation error shoots back up. This encouraged me to lower the learning rate slightly. I found 0.05 to be a good balance; even though there is still some overshoot, it retains a relatively fast convergence rate compared to smaller learning rates.

```
epoch: 14
training loss: 0.21323843662272415
training error: 0.06282
validation loss: 0.26474000048925406
validation error: 0.0749
test loss: 0.28787428125176373
test error: 0.0794
```



