



P1

Monday, March 22, 2021 2:14 PM



Predicted Velocity Graph



Ground Truth Velocity Graph

Model:

I used Pytorch's LSTMCell module for both train and test in order to control the model parameters more easily. I also defined a nn.Linear layer to transform the output so that its last dimension is the same as the inputs', not the hidden layers'.

For training, since the hidden states h_x and c_x should be of dimensions (batch_size, self.hidden_size), I defined a helper function init_hidden_state to initialize them to zero. Using a for loop through each timestep, I called self.lstmCell and appended the output at each timestep to construct a final output (the predictions).

For testing, there is only one initial 'x' predicting multiple outputs; this is a One-to-Many structure. In my for loop through the timesteps, I propagate the output of each lstmCell into the next lstmCell's input.


Loss Function: Since this is essentially a regression problem, I used Mean-Squared Error loss to evaluate the loss between the predicted velocities and the ground truth.

Test L1 error: 1.0259

Test L2 error: 0.0036

Hyperparameters:

num_epochs = 5
lr = 0.0025 (no learning rate decay/scheduler)
num_layers = 2
dropout = 0.1



When I initially experimented with 20 epochs with a learning rate of 0.005, I observed a training loss graph that drastically diverged after 3 epochs. This meant the learning rate was too high. In addition, the L1 error was quite high (1.98).

Using a smaller learning rate of 0.0025 helped to prevent overshooting, but still ensured quick convergence. Increasing the num_layers to four did not significantly boost the accuracy, and also took longer to run because the recurrent depth is greater.